

# VE492 Final Recitation Class

Yunpeng Jiang, Zhengjie Ji

UMJI

*{jyp9961, jizhengjie}@sjtu.edu.cn*

July 29, 2021

# Table of Contents

- 1 Probability
- 2 Bayes Nets
- 3 Hidden Markov Models
- 4 Introduction to ML
- 5 Discriminative Learning

# Probability - Outline


- Random Variables
- Joint and Marginal Distributions
- Conditional Distribution
- Product Rule, Chain Rule, Bayes' Rule
- Inference
- (Conditional) Independence

# Probability - Background knowledge

- $X, Y$  independent if and only if  $\forall x, y : P(x, y) = P(x)P(y)$
- $X, Y$  are conditionally independent given  $Z$  if and only if:  
 $X \perp\!\!\!\perp Y|Z, \quad \forall x, y, z : P(x, y|z) = P(x|z)P(y|z)$
- Conditional Probability:  $P(x|y) = P(x, y)/P(y)$
- Product rule:  $P(x, y) = P(x|y)P(y)$
- Chain rule:  $P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i|X_1, \dots, X_{n-1})$
- Sum rule (marginalization):  $p(X) = \sum_y P(X, y)$
- Variant of sum rule:  $p(X) = \sum_y P(X|y)P(y)$
- Bayes rule:  $P(y|x) = p(x|y)p(y)/P(x)$

# Quick Example: Bayes Rule

- $P(Z) = 0.02$  (zebra in 2% of images)
- $P(O|Z) = 0.8$  (true positive)
- $P(O|Z) = 0.1$  (false positive)
- We want to calculate:  $p(Z|O)$
- Apply Bayes rule!




# Bayes Nets - Outline

- Representation
- Conditional Independence
  - D-separation
  - Active / Inactive Paths
- Probabilistic Inference
  - Enumeration
  - Variable elimination
  - Probabilistic inference
  - Sampling

# Bayes Nets - Sample Questions

- How to get formula of joint distribution from BN graphs?
- How to count the degree of freedom of BN graphs?
- How to run variable elimination? What is the best ordering for VE? What is the largest generated factor? What is the cutset?
- What are the difference of the four sampling methods? What is the time complexity?

## Example: Small Bayes Net



- Provide the formula of the joint distribution over all the variables given by the Bayes net.
- Provide the number of degrees of freedom of the BN.
- Run VE to compute to compute  $P(A|H = h)$ . Provide the list of the sizes (i.e., number of variables) of the factors obtained at the end of each iteration in VE.



# Hidden Markov Models - Outline


- Markov Models and Hidden Markov Models
- Forward algorithm
- Viterbi algorithm
- Particle filtering
- Dynamic Bayesian Network

# Hidden Markov Models - Sample Questions

- How to compute stationary distribution?
- How to do filtering / prediction / smoothing / explanation?
- What is particle filtering?
- How to distinguish Dynamic Bayes Nets and Bayes Nets?

# HMM Terminology

Time instants	$t$ in $\{1, 2 \dots T\}$
Hidden States / States / Emitters	$X_t$
Outputs / Emissions / Observations / Visible States	$O_t$
All possible states / states set	$X_t$ in $\{1, 2 \dots N\}$
All possible emissions / emissions set	$O_t$ in $\{1, 2 \dots K\}$
Initial state distribution / Initial state probabilities	$p_i$ in $q$ or $\pi_i$ in $\pi$
Transition probabilities / State transition probabilities	$a_{ij}$ in row-stochastic matrix $A$
Emission probabilities / Observation probabilities	$b_{jk}$ in row-stochastic matrix $B$




# Example - Adapt the Forward Algorithm

## Filtering Algorithm

$$P(X_{t+1}|e_{1:t+1}) = \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(x_t|e_{1:t}) P(X_{t+1}|x_t)$$

- Adapt the forward algorithm to this variant of HMM.
- Step 1. Predict step
- Step 2. Update



# Hidden Markov Models - Background knowledge

## Forward Algorithm (rewrite)

$$\alpha_t(i) = \left[ \sum_{j=1}^N \alpha_{t-1}(j) a_{ji} \right] b_i(O_t)$$

## Viterbi Algorithm (rewrite)

$$\delta_t(i) = \max_{j \in \{1, \dots, N\}} [\delta_{t-1}(j) a_{ji} b_i(O_t)]$$

- $a_{ij}$ : state transition probabilities,  $b_{jk}$ : emission probabilities
- Forward Algorithm can be used to predict the current state given all of the current and past evidence.
- Viterbi algorithm can be used to calculate the most likely state sequence, namely  $\operatorname{argmax}_{X_{1:T}} P(X_{1:T} | O_{1:T}, \lambda)$ , where  $\lambda = \{A, B, \pi\}$ .

Please refer to the lecture slides if this leads to any confusion!



# Example: Viterbi Algorithm

**A:  $a_{ji} = P(X_{t+1} = j | X_t = i)$**

$X_t   X_{t+1}$	A	B	H	S
A	0.6	0.1	0.1	0.2
B	0.0	0.3	0.2	0.5
H	0.8	0.1	0.1	0.1
S	0.2	0.0	0.1	0.7

**B:  $b_{ik} = P(O_t = k | X_t = i)$**

$X_t   O_t$	p	e	b	l
A	0.6	0.2	0.1	0.1
B	0.1	0.4	0.1	0.4
H	0.0	0.0	0.7	0.3
S	0.0	0.0	0.1	0.9

**$\pi = P(X_1 = i)$ :**

A	B	H	S
0.5	0.0	0.0	0.5

**Observations :**

**$o_{1:4} = \{ b, p, l, e \}$**

**Find:**

**Most likely hidden state**

**sequence:  $X_{1:4}^*$**

## Naive Bayes model for classification

- Naive Bayes assumes all features are independent effects of the label.
- Naive Bayes for text:  $P(Y, W_1, \dots, W_n) = P(Y) \prod_i P(W_i | Y)$ , where  $W_i$  is the word at position  $i$ ,  $Y \in \{\text{spam}, \text{ham}\}$ .



## Maximum likelihood estimation

- Given the observed set  $D$ , find  $\theta$  to maximize the probability of  $D$
- $\theta = \underset{\theta}{\operatorname{argmax}} P(D|\theta)$
- Set the derivative of  $P(D|\theta)$  with respect to  $\theta$  to zero, and solve for  $\theta$ .

## Laplace smoothing

- pretend that we have seen every outcome  $k$  extra times.
- $P_{Lap,k} = \frac{c(x)+k}{N+k|X|}$

# Discriminative Learning

## Linear classifiers

feature values(inputs), weights (learned), activation (sum,  
 $\text{activation}_w(x) = \sum_i w_i \phi_i(x)$ )

## Binary perceptron learning process

- start with weights=0
- for each training instance  $(x, y^*)$ : classify with current weights, no change if correct else adjust the weight vector by adding or subtracting the feature vector (subtract if  $y^*=-1$ ).

# Discriminative Learning

## Multiclass perceptron learning process

- start with weights=0
- pick up training examples one by one
- predict with current weights  $\hat{y} = \operatorname{argmax}_y (w_y \cdot \phi(x))$ , no change if correct. Otherwise, lower score of wrong answer and raise score of right answer  $w_{\hat{y}} = w_{\hat{y}} - \phi(x)$ ,  $w_{y*} = w_{y*} + \phi(x)$

## Probabilistic Perceptron

- softmax function

# Discriminative Learning

## Learning by gradient descent

initialize  $w$  (e.g., randomly)

repeat for  $K$  iterations:

for each example  $(x_i, y_i)$ :

compute gradient  $\Delta_i = -\nabla_w \log P_w(y_i|x_i)$

compute gradient  $\nabla_w \mathcal{L} = \sum_i \Delta_i$

$w \leftarrow w - \alpha \nabla_w \mathcal{L}$

$$\frac{d}{dw_y} \log P_w(y_i|x_i) = x_i(I(y = y_i) - P(y|x_i))$$

- ❖  $\alpha$ : learning rate — hyperparameter that needs to be chosen carefully
- ❖ How? Try multiple choices
  - ❖ Crude rule of thumb: update should change  $w$  by about 0.1-1%

# The End